

ServiceGuard AI

Product Definition & Vision

1. Product statement

ServiceGuard AI is an AI-powered SaaS/service operations control tower that continuously analyzes application telemetry, predicts SLA/service risk, retrieves approved operational runbooks, and provides evidence-backed recommendations to operations engineers.

2. Problem

- Operational teams monitor multiple services across customers and infrastructure.
- Signals are distributed across monitoring systems, incident tools and documentation.
- Engineers spend time correlating metrics with historical incidents and runbooks.
- Early degradation is difficult to prioritize before an SLA breach.
- LLMs alone are unreliable because they do not inherently know current telemetry or approved procedures.

3. Product objective

Reduce mean time to detect, understand and respond to service degradation while preserving human control and auditability.

4. Core workflow

```
Service telemetry
↓
Anomaly / SLA detection
↓
Incident
↓
Context collection
↓
ML risk prediction
↓
Runbook / SOP retrieval
↓
AI Agent synthesis
↓
Recommended action
↓
Engineer approval
↓
Resolution + audit
```

5. Success metrics

- Incident detection latency.
- Percentage of incidents with complete evidence.
- Recommendation acceptance rate.
- False-negative rate of critical-risk prediction.
- Mean time to acknowledge and resolve incidents.
- RAG retrieval relevance.
- AI response groundedness.
- API availability and latency.

6. Product boundaries

- MVP is decision support, not autonomous production remediation.
- ML predicts operational risk; it does not replace deterministic alert rules.
- RAG supplies approved procedures; it does not execute them.
- Production changes require explicit authorization.